MODEL SELECTION, TRAINING & EVALUATION REPORT

Machine Learning & Deep Learning Prediction Project

Dual-Branch Pediatric Pneumonia Detection — Phase 2

Team

1. **Kassouar Fatima** 222238233310
   Project Manager · CSV Branch · ML .Fusion · Report
2. **Bouhmidi Amina Maroua** 222238192905
   Data Engineer · X-Ray Branch
3. **Labani Nabila Nour Elhouda**
   DL Engineer · X-Ray Branch
4. **Miloudi Meroua Amira** 222238200409
   Developer · Business Side

University of Saida, Algeria

Supervisor: Dr. Abderrahmane Khiat
Medical Advisor: Dr. Aimer Mohammed Djamel Eddine

Project Title:
AI-Powered Early Pediatric Pneumonia Detection for Kids & EMR Integration in Algeria

March 3, 2026
Phase 2 — Model Training & Evaluation

Prediction Target: Pneumonia (Yes/No) — Binary classification from clinical vitals (CSV) and chest X-ray images

Task Type: ☑ **Binary Classification** ☐ Multi-class ☐ Regression

Preprocessing Phase: ☑ Completed & Validated

GitHub : https://github.com/labaninabila193-code/AI-Pediatric-Pneumonia-Detection

GitHub Result : GitHub: https://github.com/labaninabila193-code/AI-Pediatric-Pneumonia-Detection | Models on Google Drive (shared link in README).

🎯 SECTION 1 — Task Framing & Model Selection Criteria

1.1 Problem Framing Summary

Before selecting any model, the team clearly defined what was being predicted and what a good prediction looks like clinically. This project uses **two parallel prediction branches**: the CSV/Clinical Branch and the X-Ray CNN Branch.

What exactly is the model predicting? (output variable, format, scale)	**CSV Branch:** Binary probability score ∈ [0,1] per patient clinical record based on 12 clinical features (Fatigue, Fever, Temperature, Cough, Breathing Difficulty, etc.). Final label: PNEUMONIA (1) if score > 0.5, NORMAL (0) otherwise. Best model: Gradient Boosting Classifier. **X-Ray Branch:** Binary probability score ∈ [0,1] per chest X-ray image (224×224 RGB). Final label: PNEUMONIA (1) if score > 0.26 (optimised threshold), NORMAL (0) otherwise. Best model: DenseNet121 (ImageNet transfer learning).
What does a correct prediction look like?	**CSV Branch:** True Positive: patient with pneumonia flagged as PNEUMONIA → immediate antibiotic prescription. True Negative: healthy patient correctly cleared as NORMAL → no unnecessary treatment. **X-Ray Branch:** True Positive: chest X-ray with consolidation correctly classified as PNEUMONIA. Grad-CAM heatmap highlights lung consolidation region. True Negative: clear X-ray correctly classified as NORMAL.
What does a wrong prediction cost? (clinical / business impact)	**Both Branches — False Negative:** delayed antibiotic therapy → risk of sepsis, respiratory failure, or death in a child. Clinically catastrophic. **Both Branches — False Positive:** unnecessary antibiotic course, parental anxiety, possible over-radiation — manageable and correctable by physician.
Which error is more costly — False Negative or False Positive? Justify:	**FALSE NEGATIVE is far more costly — for both branches.** Paediatric pneumonia can deteriorate into life-threatening sepsis within hours. Medical standards universally prioritise sensitivity ≥ 95%. This justifies: (1) lowering CNN threshold from 0.5 → 0.26 to achieve ≥95% sensitivity; (2) applying class_weight="balanced" in all CSV models.

1.2 Constraints & Requirements

Constraint	Value / Choice	Justification
Acceptable inference time	☑ Near-real (~10s) — X-Ray ☑ Real-time (<1s) — CSV	Hospital radiologist workflow tolerates ~10 s/image. CSV clinical prediction is nearly instantaneous.
Deployment environment	☑ Hospital GPU server (X-Ray) ☑ Cloud / API (CSV)	Models saved as .keras and .pkl for hospital GPU server and API deployment in Algeria.
Model interpretability needed?	☑ Yes — both branches	Clinicians require Grad-CAM heatmaps (X-Ray) and SHAP / feature importance (CSV). Builds clinical trust.
Maximum model size (MB)	~30 MB (DenseNet121) <1 MB (CSV .pkl files)	DenseNet121 (30 MB) is lighter than VGG16 (528 MB) while outperforming it.
Minimum acceptable sensitivity	≥ 95% (X-Ray) — Achieved: 95.01% ☑ ≥ 80% (CSV) — Achieved: 85.71%	Medical standard: ≥ 95% sensitivity ensures no more than 1 in 20 pneumonia cases is missed.
Minimum acceptable specificity	≥ 80% — X-Ray: 92.44% ☑ 100% — CSV	Below 80% generates unacceptably high false alarm burden on radiology departments.
Framework constraint	☑ TensorFlow/Keras (X-Ray) ☑ Scikit-learn (CSV)	TensorFlow/Keras for deep learning; Scikit-learn for classical ML branch.

SECTION 2 — Model Candidates & Selection Rationale

⚠ Instructor Note — Common Mistakes to Watch For

• Students must always start with a baseline model before complex ones.
• Model choice must be justified by the task type — not just "I used ResNet because it is famous".
• Transfer learning requires justification of why the source domain is relevant.
• Custom architectures must explain every design decision.

2.1 Baseline Model (Mandatory)

A baseline is mandatory. It sets the performance floor against which all other models are judged. One baseline per branch was established before any complex model was trained.

CSV / Clinical Branch

Parameter	Value	Notes
Baseline model type	Logistic Regression	Simple linear classifier on 12 clinical features; n=127 patients.
Baseline validation accuracy	87.50%	Evaluated on 32 paediatric test patients (25% hold-out split).
Baseline AUC-ROC	0.9206	Surprisingly high — clinical features carry strong discriminatory signal.
Baseline F1-Score (macro)	0.8462	Acceptable but below Gradient Boost F1 = 0.9231.
Why this baseline was chosen	Most interpretable binary classifier	Reveals whether the problem is linearly separable. AUC = 0.92 confirmed the task is well-structured.

What did the CSV baseline reveal about the difficulty of the task?

Clinical features carry very strong discriminatory signal — AUC = 0.92 on only 127 patients confirms the task is somewhat linearly separable. However, recall of 78.57% at default threshold is insufficient for clinical use (target ≥ 80%). This justifies pursuing ensemble models with threshold optimisation. The gap between baseline (87.5%) and best CSV model (93.75%) validated pursuing Gradient Boosting.

X-Ray CNN Branch

Parameter	Value	Notes
Baseline model type	Simple Custom CNN (3 Conv layers)	Mandatory baseline to prove why transfer learning is needed.
Baseline validation accuracy	~72%	Insufficient — training from scratch on 4,119 images does not work.
Baseline AUC-ROC	~0.75	Low — confirms deep feature extraction requires pretrained weights.
Baseline F1-Score	~0.73	Insufficient for clinical deployment.
Why this baseline was chosen	Establishes performance floor before transfer learning	Demonstrates that a 3-layer custom CNN cannot extract sufficient chest X-ray features from ~4K training images.

What did the X-Ray baseline reveal about the difficulty of the task?

Training from scratch is infeasible at ~4K images. The custom CNN baseline AUC ≈ 0.75 confirms that chest X-ray feature extraction requires deep hierarchical representations impossible to learn without pretrained initialisation. This directly validated the transfer learning approach. The baseline also revealed class imbalance: ~73% PNEUMONIA / 27% NORMAL, requiring class_weight="balanced".

2.2 Candidate Models Considered — CSV Branch

The CSV branch involves 12 structured clinical features on 127 patients — a small tabular dataset. Deep neural networks overfit severely on n<200 tabular data. All 7 classical ML models were considered and compared.

Model Name	Architecture	Considered?	Selected?	Reason for Selection or Rejection
Logistic Regression	Linear / Classical ML	✅ Yes	✅ BASELINE	Mandatory baseline. AUC=0.9206, Accuracy=87.5%. Insufficient sensitivity (78.57%) for clinical use.
Random Forest	Ensemble / Classical ML	✅ Yes	✅ Yes (backup)	AUC=0.9782, Accuracy=93.75%. Selected as backup model. Handles non-linear interactions.
SVM	Kernel-based	✅ Yes	✗ No	AUC=0.9405, Accuracy=87.5%. Below Gradient Boost. Kernel trick adds no value for 12 features. Rejected.

Model Name	Architecture	Considered?	Selected?	Reason for Selection or Rejection
KNN	Instance-based	☑ Yes	✗ No	AUC=0.9762 but CV variance ±0.068 — unreliable on n=127. Sensitive to scaling. Rejected.
Naive Bayes	Probabilistic	☑ Yes	✗ No	AUC=0.9841 but independence assumption invalid — Fatigue correlates with Fever. Rejected.
MLP Neural Network	Neural Network	☑ Yes	✗ No (WORST)	AUC=0.6905, CV AUC=0.5834±0.174 — nearly random. Textbook failure of DL on small tabular data. Rejected.
XGBoost / Gradient Boost ★ FINAL	Gradient Boosting	☑ Yes	★ BEST	AUC=0.9960, Accuracy=93.75%, F1=0.9231, CV AUC=0.9939±0.009. Zero false positives. Selected.

2.2 (continued) — Candidate Models Considered — X-Ray CNN Branch

Model Name	Architecture	Considered?	Selected?	Reason for Selection or Rejection
Simple Custom CNN	Deep Learning — Custom	☑ Yes	☑ BASELINE	AUC~0.75. Mandatory baseline. Proves transfer learning is needed.
ResNet50	Deep Learning — Transfer	☑ Yes	✗ No (REJECTED)	AUC=0.8802. At 95% sensitivity threshold, specificity collapses to 31.51% — 68.5% false alarms. Clinically unacceptable.
VGG16	Deep Learning — Transfer	☑ Yes	✗ Not final	AUC=0.9644, Sensitivity=95.32% but Specificity=80.67% (marginal). Also 528 MB — too large for hospital deployment.
EfficientNet	Deep Learning — Transfer	☑ Yes	✗ No	Not evaluated in detail. AUC estimate ~0.96 but higher compute cost vs DenseNet121. Not pursued given DenseNet121 met all targets.

Model Name	Architecture	Considered?	Selected?	Reason for Selection or Rejection
DenseNet121 ★ FINAL	Deep Learning — Transfer	✅ Yes	★ BEST	AUC=0.9810, Accuracy=94.31%, Sensitivity=95.01% ✅, Specificity=92.44% ✅. Only model meeting ALL clinical thresholds. External validation: Sensitivity=97.21%.

2.3 Transfer Learning Configuration — X-Ray CNN Branch

This section applies to the X-Ray CNN Branch only. The CSV Branch uses classical ML — no pretrained models.

Pretrained weights source	ImageNet (1.2 million images, 1000 classes — Keras official weights)
Relevance of source domain to your task	ImageNet low-level features (edges, textures, gradient detectors) are universal visual primitives that transfer directly to radiological images. Consolidation patterns, opacity boundaries, and lung architecture involve the same low-level statistical regularities as natural image edges and textures. Studies consistently confirm ImageNet-pretrained CNNs outperform random initialisation for chest X-ray tasks at small dataset scales.
Layers frozen during Phase 1	All 427 layers of DenseNet121 base model frozen. Only the custom classification head was trained GlobalAveragePooling → Dense(512, ReLU) → Dropout(0.5) → Dense(1)
Layers fine-tuned during Phase 2	Not performed. Frozen backbone already achieved all clinical targets. Future work: unfreeze last 2 dense blocks at LR=0.00001 for potential AUC improvement to ~99%.
Custom classification head architecture	GlobalAveragePooling → Dense(512, ReLU) → Dropout(0.5) → Dense(1)
Reason for this head design	GlobalAveragePooling replaces Flatten to dramatically reduce parameters and prevent overfitting. BatchNorm stabilises gradient magnitudes. Dropout(0.5) provides aggressive regularisation on ~4K images. Single sigmoid output with binary cross-entropy is standard for binary classification.

Input size compatibility confirmed?	☑ YES — DenseNet121, ResNet50, VGG16 all expect 224×224×3 RGB. Images resized via ImageDataGenerator with rescale=1./255.

Why is the pretrained model's source domain (ImageNet) relevant or adaptable to chest X-rays?

Although ImageNet contains natural photographs, the convolutional feature detectors learned at early layers (Gabor-like filters, colour gradients, texture detectors) are domain-agnostic and transfer directly to medical images. The consolidation patterns and opacity boundaries detected in chest X-rays involve the same low-level statistical regularities as natural image edges. Even without fine-tuning the backbone, the rich ImageNet feature representations are sufficient for a trained classification head to learn pneumonia-specific patterns from the final feature maps. This is why AUC = 0.981 was achieved with a frozen backbone on only 4,119 training images.

⚙ SECTION 3 — Training Configuration

3.1 Hardware & Environment

Hardware used (GPU/CPU/TPU)	Framework & Version	Training Duration (approx.)
NVIDIA T4 GPU (Google Colab — 16 GB VRAM) — X-Ray CPU (Google Colab) — CSV	TensorFlow 2.x / Keras 3.x (X-Ray) scikit-learn 1.6.1 (CSV) Python 3.10–3.12	DenseNet121: ~57.97 min | ResNet50: ~65 min VGG16: ~80 min | CSV all 7 models: < 1 min
Environment: Google Colab (cloud Jupyter) with Google Drive mounted. GPU runtime T4. Outputs saved to /content/drive/MyDrive/Pneumonia_Project/	Random Seed: Seed = 42 in ALL model configs. CSV: random_state=42 in all sklearn models and train_test_split. CNN: tf.random.set_seed(42).	Total training time (fit + 5-fold CV, all 7 models): 1.24 s

3.2 Hyperparameter Configuration

X-Ray CNN Branch

Hyperparameter	Value Used	Search Range Tried	Final Chosen	Justification
Initial learning rate	0.0001	0.001 / 0.0001	0.0001	LR=0.001 caused instability.

Hyperparameter	Value Used	Search Range Tried	Final Chosen	Justification
				LR=0.0001 gave smooth convergence.
Learning rate schedule	ReduceLROnPlateau	Fixed / ReduceLR	ReduceLROnPlateau	Reduces LR by 0.5 when val_loss plateaus 3 epochs. Smooths later training.
Optimizer	Adam	Adam only	Adam	Standard for transfer learning. Adaptive per-parameter LR.
Optimizer beta1	0.9 (default)	Default	0.9	Well-validated default for image classification.
Batch size	32	16 / 32	32	Maximises T4 GPU utilisation while maintaining gradient noise.
Number of epochs (total)	30 max → stopped ep.16 (best: ep.9)	20 / 30	30 (EarlyStopping)	EarlyStopping at patience=7. Best model at epoch 9.
Early stopping patience	7 epochs	5 / 7 / 10	7	Allows recovery from temporary plateaus before stopping.
Dropout rate	0.5	0.3 / 0.5	0.5	Aggressive dropout (50%) counteracts overfitting on ~4K images.
L1/L2 regularisation	None explicit	N/A	None	BatchNorm + Dropout provide

Hyperparameter	Value Used	Search Range Tried	Final Chosen	Justification
				sufficient implicit regularisation.
Weight initialisation	ImageNet (base) + He Normal (head)	N/A	ImageNet	Transfer learning starts from convergent weights.
Loss function	☑ Binary Cross-Entropy	Binary CE	Binary CE	Standard for binary sigmoid output.
Class weights applied?	☑ Yes — class_weight="balanced"	balanced	Balanced	Training set is ~73% PNEUMONIA. Balanced weights correct majority-class bias.

CSV / Clinical Branch

Hyperparameter	Gradient Boost	Random Forest	Logistic Reg.	Notes
n_estimators	100	100	N/A	Chosen via cross-validation as optimal bias-variance point.
learning_rate	0.1 (default)	N/A	C=1.0	GBM default 0.1 proved stable on n=127.
max_depth	3 (shallow)	None (full)	N/A	Shallow trees limit capacity, reduce overfitting.
seed /random_state	42	42	42	Ensures full reproducibility.
class_weight	balanced	balanced	balanced	Corrects for class imbalance in 127 patients.

Hyperparameter	Gradient Boost	Random Forest	Logistic Reg.	Notes
Train/test split	75% / 25% stratified	75% / 25% stratified	75% / 25% stratified	Stratification preserves class proportions.

3.3 Training Strategy

X-Ray CNN Branch

Strategy Element	Applied?	Details
Phase 1: Frozen backbone (transfer learning)	☑ Yes	All 427 DenseNet121 layers frozen. Only classification head trained. Epochs 1–16. LR: 0.0001.
Phase 2: Fine-tuning unfrozen layers	✗ No	Not performed. Frozen backbone achieved clinical targets. Future: unfreeze last 2 dense blocks at LR=0.00001.
Early stopping	☑ Yes	Monitor: val_loss | Patience: 7 | Mode: min | Restore best weights: True. Triggered epoch 16, best epoch 9.
Model checkpoint (save best)	☑ Yes	Monitor: val_loss | save_best_only: True | File: densenet121_best_model.keras
Learning rate warm-up	✗ No	Starting LR=0.0001 is already conservative enough.
Gradient clipping	✗ No	Not needed. DenseNet121 dense connections stabilise gradients inherently.
Mixed precision training	✗ No	Not required. T4 GPU with float32 provides sufficient speed.
Class weight integration	☑ Yes	class_weight="balanced" passed to model.fit(). Corrects 73%/27% imbalance.
Data augmentation	☑ Yes	Rotation ±20°, shift 0.2, zoom 0.2, horizontal flip. Applied ONLY to training generator.

📊 CSV / Clinical Branch

Strategy Element	Applied?	Details
Cross-validation	☑ Yes	5-Fold CV, RepeatedKFold (3×5=15 folds), Bootstrap (1000 samples), LOOCV used to validate model stability.
Stratified train/test split	☑ Yes	75%/25% stratified split. Preserves class proportions in both sets.

Strategy Element	Applied?	Details
Threshold optimisation (CSV)	☑ Yes	Default threshold=0.5 retained. Gradient Boost already achieves 85.71% sensitivity / 100% specificity at 0.5.

3.4 Callbacks & Monitoring

Callbacks used (X-Ray CNN Branch):

- ModelCheckpoint — filepath: densenet121_best_model.keras | monitor: val_loss | save_best_only: True
- EarlyStopping — monitor: val_loss | patience: 7 | mode: min | restore_best_weights: True
- ReduceLROnPlateau — monitor: val_loss | factor: 0.5 | patience: 3 | min_lr: 1e-7
- TensorBoard — log_dir: /content/drive/MyDrive/Pneumonia_Project/Models/logs | histogram_freq: 1

CSV Branch: No callbacks required. sklearn GBM.fit() trains sequentially. Cross-validation strategies used instead.

model.compile() and model.fit() code snippet (X-Ray CNN Branch):

```
model.compile(
    optimizer=Adam(learning_rate=0.0001),
    loss='binary_crossentropy',
    metrics=['accuracy', AUC(name='auc')])
history = model.fit(
    train_generator, epochs=30, validation_data=val_generator,
    callbacks=[checkpoint, early_stopping, reduce_lr,
tensorboard],
    class_weight='balanced')
```

SECTION 4 — Learning Curves & Training Behaviour

Required Plots

- Training loss vs Validation loss (over all epochs) — both branches
- Training accuracy vs Validation accuracy (over all epochs)
- Learning rate schedule plot (ReduceLROnPlateau applied to CNN branch)
- For multi-phase training: mark the phase transition clearly on the plot

Plots generated in training notebooks → saved to /content/drive/MyDrive/Pneumonia_Project/Models/figures/ CNN

4.1 Loss Curve Analysis

X-Ray CNN Branch

Epoch at which validation loss starts stabilising	Epoch ~6–7. Validation loss showed consistent improvement until epoch 9 (best checkpoint), then plateaued and slightly increased.
Lowest training loss achieved	~0.12–0.15 (best at epoch 9, logged in TensorBoard and training CSV history)
Lowest validation loss achieved	Achieved at epoch 9 — this is the model saved by ModelCheckpoint.
Gap between training and validation loss at final epoch	Small gap (~0.05–0.10), indicating mild but controlled overfitting.
Was early stopping triggered? At which epoch?	YES — triggered at epoch 16 (patience = 7 from best epoch 9). Best weights from epoch 9 restored.
Learning rate at which best performance was achieved	0.0001 (initial). ReduceLROnPlateau reduced to ~0.00005 around epoch 12–14.

Describe the shape of your loss curves (X-Ray Branch). What story do they tell?

Training loss (DenseNet121): Steadily decreasing from epoch 1 through epoch 9, then continuing slowly through epoch 16. Validation loss mirrored training loss until epoch 9 (best checkpoint), then plateaued and slightly increased — a classic sign of memorisation beginning.

The story: DenseNet121 learned rapidly in the first 5 epochs (large gradient updates on the new classification head, with a frozen ImageNet backbone providing stable feature maps). The curve then refined gradually. The mild divergence after epoch 9 never becomes alarming, confirming the regularisation strategy (Dropout 0.5 + BatchNorm + EarlyStopping) worked well.

ReduceLROnPlateau smoothed later epochs by preventing oscillation near the optimum.

CSV / Clinical Branch

Describe the shape of your loss curves (CSV Branch). What story do they tell?

CSV Gradient Boosting: Training loss decreases monotonically as n_estimators increases (each tree corrects errors of the previous). Validation CV AUC stabilised around estimator 60–80 and remained at AUC = 0.9939±0.009.

The story: Gradient Boosting on clinical features converges rapidly because the 12 features carry a strong, clear diagnostic signal. Mild overfitting begins at n_estimators > 150, confirming n_estimators = 100 as the optimal choice. The very small CV variance (±0.009) demonstrates the model learns generalisable patterns despite only 127 patients.

4.2 Overfitting / Underfitting Diagnosis

X-Ray CNN Branch

Diagnosis Pattern	Observed?	Corrective Action Taken
Overfitting (train loss ↓, val loss ↑)	☑ MILD	Dropout (0.5), BatchNormalization, EarlyStopping (patience=7, restores best weights), ReduceLROnPlateau, Data Augmentation.
Underfitting (both losses high/stagnant)	✗ No	ImageNet weights provide excellent initialisation — no flat early-training loss.
High variance (val loss oscillates)	✗ No	ReduceLROnPlateau smoothed oscillations by reducing LR when plateau detected.
Vanishing gradient (loss stuck early)	✗ No	DenseNet121's dense connections (skip connections to all previous layers) inherently prevent vanishing gradients.
Divergence (loss increases)	✗ No	Adam with conservative LR=0.0001 prevented divergence throughout training.

WHY did mild overfitting occur? (X-Ray Branch) — and what was done to address it?

ROOT CAUSE 1 — Dataset size mismatch with model capacity: DenseNet121 backbone has ~7 million parameters. The training set has only ~4,119 images. Even with a frozen backbone, the trainable classification head has sufficient capacity to partially memorise training-set-specific feature distributions from the final DenseNet feature maps. This capacity–data imbalance is the PRIMARY driver of mild overfitting.

ROOT CAUSE 2 — Class imbalance (73% PNEUMONIA): Even with class_weight="balanced", the model sees 2.7× more pneumonia examples per epoch. The decision boundary is learned predominantly from the majority class, causing the head to slightly memorise majority-class-specific patterns. This manifests as training loss continuing to drop after epoch 9 while validation loss plateaus.

ROOT CAUSE 3 — Frozen backbone bottleneck: The frozen backbone applies the same ImageNet feature maps to both augmented training images and non-augmented validation images. The head memorises the specific mapping between augmented-train feature activations and labels. Minor distribution shift between augmented train and clean validation explains the small train–val gap.

Actions taken: (1) Dropout(0.5) — 50% of head neurons disabled per forward pass. (2) BatchNormalization — re-centres activations, acts as additional regularisation. (3) EarlyStopping(patience=7, restore_best_weights=True) — deploys epoch-9 weights before overfitting grows. (4) Data Augmentation — increases training diversity. (5) ReduceLROnPlateau — smaller gradient steps prevent jumping over the validation optimum.

📊 CSV / Clinical Branch

Diagnosis Pattern	Observed?	Corrective Action Taken
Overfitting (train AUC >> CV AUC)	✅ MILD at n_estimators > 150	n_estimators=100 chosen as sweet spot. class_weight="balanced". Shallow trees (max_depth=3).
Underfitting	✗ No	Gradient Boosting naturally avoids underfitting via iterative error correction.
High variance across CV folds	✗ No	CV AUC = 0.9939±0.009 — extremely low variance. Model is stable.

WHY did mild overfitting occur? (CSV Branch) — and what was done to address it?

ROOT CAUSE 1 — Extremely small sample (n=127): With only 95 training patients, Gradient Boosting with 100 trees × depth-3 has theoretical capacity for up to 800 leaf nodes — vastly exceeding the training set. The model memorises specific patient-symptom co-occurrence patterns rather than generalising. This is why AUC=0.9960 on training but some variance is expected in external deployment.

ROOT CAUSE 2 — Dominant single feature (Fatigue = 64.8%): When one feature dominates this strongly, the model develops high confidence on training samples where Fatigue is clearly present or absent. Borderline cases where Fatigue is moderate are fitted with a precise decision boundary that may not generalise. Shallow trees (max_depth=3) partially mitigate this.

Actions taken: (1) max_depth=3 — shallow trees limit individual tree complexity. (2) class_weight="balanced" — prevents majority-class bias. (3) n_estimators=100 — chosen via 5-fold CV as optimal. (4) Multiple validation strategies (Bootstrap, LOOCV, RepeatedKFold) confirmed stability. CV AUC = 0.9939±0.009 confirms mild overfitting is controlled.

🔧 SECTION 5 — Hyperparameter Tuning

5.1 Tuning Method

Method	Used?	Tool / Library	Notes
Manual search	☑ Yes	Colab / manual comparison	Primary method. Tested LR ∈ {0.001, 0.0001}, dropout ∈ {0.3, 0.5}, 3 CNN architectures, 7 CSV models.
Grid search	✗ No	—	Full grid would require 20+ Colab sessions. Impractical within project timeframe.
Random search	✗ No	—	Manual search sufficient given small hyperparameter space.
Bayesian optimisation	✗ No	—	Future: Keras Tuner (CNN head) and Optuna (CSV). Not applied due to time constraints.
Learning rate finder	✗ No	—	Conservative LR=0.0001 was set from literature. Finder not needed.
Decision threshold optimisation ★	☑ Yes (CNN)	Custom loop 0.01→0.99	CRITICAL for clinical use. Found optimal threshold=0.26 achieving 95.01% sensitivity / 92.44% specificity.
Cross-validation tuning	☑ Yes (CSV)	scikit-learn CV strategies	5-Fold, RepeatedKFold, Bootstrap, LOOCV used to validate n_estimators=100.

5.2 Top Experiments Compared

X-Ray CNN Branch

Run #	LR	Batch	Dropout	Epochs	Val AUC	Test AUC	Test Sens.	Notes
1	0.001	32	0.5	12	0.921	0.880	84.87%	ResNet50 — unstable, specificity collapses to 31.51% at 95% sensitivity. Rejected.
2	0.0001	32	0.5	30(→16)	0.973	0.964	95.32%	VGG16 — good AUC but specificity=80.67% (marginal), 528 MB too large.
3	0.0001	32	0.3	30(→14)	0.968	0.972	92.5%	DenseNet121 dropout=0.3 — more overfitting, below 95% sensitivity target.
4	0.0001	32	0.5	30(→16)	0.978	0.981	90.02%	DenseNet121 threshold=0.5 (default) — AUC excellent but sensitivity below 95% target.
5 ★	0.0001	32	0.5	→ep.9	0.981	0.981	95.01% ✅	DenseNet121 threshold=0.26 ★ BEST. Both clinical targets met.

📊 CSV / Clinical Branch

Model	Test Acc	AUC	F1	CV AUC	Notes
Logistic Regression (Baseline)	87.50%	0.9206	0.8462	—	Insufficient sensitivity (78.57%).
SVM	87.50%	0.9405	0.8571	—	Below Gradient Boost. Rejected.
KNN	87.50%	0.9762	0.8571	±0.068	High variance. Unreliable on n=127. Rejected.
Naive Bayes	87.50%	0.9841	0.8571	—	Independence assumption violated. Rejected.

Model	Test Acc	AUC	F1	CV AUC	Notes
MLP Neural Network	75.00%	0.6905	0.7143	±0.174	WORST. Nearly random. DL fails on n=127 tabular. Rejected.
Random Forest (Backup)	93.75%	0.9782	0.9231	—	Good backup model. Selected as secondary.
Gradient Boost ★ FINAL	93.75%	0.9960	0.9231	0.9939±0.009	BEST. Zero FP. Stable CV. Final CSV model.

What was the single most impactful hyperparameter change? Why did it have such an effect?

X-Ray Branch: Decision threshold optimisation (0.50 → 0.26) was the single most impactful change. Effect: Sensitivity jumped from 90.02% → 95.01% (+5pp) while overall accuracy IMPROVED (92.04% → 94.31%) because more true positives were captured. Why: DenseNet121 already learned strong feature representations (AUC=98.1%). Moving the threshold to 0.26 reclassifies borderline cases (probability ~30–50%) as PNEUMONIA — and since these are predominantly true pneumonias, both sensitivity and accuracy improve simultaneously. KEY LESSON: For medical applications, threshold selection is a CLINICAL decision, not a technical default.

CSV Branch: Switching from Logistic Regression to Gradient Boosting (+0.075 AUC). The sequential error-correction mechanism of Gradient Boosting captures non-linear Fatigue × Fever × Temperature interactions impossible for logistic regression's linear boundary.

SECTION 6 — Evaluation Metrics & Results

⚠ Critical Reminder: Accuracy Is Not Sufficient for Imbalanced Datasets

- For classification with class imbalance: always report AUC-ROC, F1, Sensitivity, Specificity.
- For regression: always report MAE, RMSE, and R^2 together.
- A model predicting the majority class (PNEUMONIA) 100% of the time achieves 73% accuracy on this dataset.
- If you only report accuracy, this section will receive 0 points.

6.1 Classification Metrics — Both Branches

Metric	Formula	DenseNet121 (Validation) [X-Ray]	DenseNet121 (Test thr=0.26) [X-Ray]	Gradient Boost (Test) [CSV]	Baseline LR (Test) [CSV]	Target
Accuracy	TP+TN/Total	~91.5%	94.31%	93.75%	87.50%	≥85%
AUC-ROC	Area under ROC	~0.979	0.9810 ★	0.9960 ★	0.9206	N/A
F1-Score (macro)	2×(P×R)/(P+R)	_	_	0.9231	0.8462	> 0.80
F1 (class 0 — Normal)	2×(P0×R0)/(P0+R0)	~0.95	0.9583	0.9444	0.8696	
F1 (class 1 — Pneumonia)	2×(P1×R1)/(P1+R1)	~0.91	0.9231	0.9231	0.8571	>0.80
Sensitivity (Recall)	TP/(TP+FN)	~94.8%	95.01% ☑	85.71%	78.57%	≥95% ☑
Specificity	TN/(TN+FP)	~93.0%	92.44% ☑	100%	100%	≥80% ☑
Precision (PPV)	TP/(TP+FP)	~96.8%	97.13%	100%	100%	—
Matthews Corr. Coeff.	(TP×TN−FP×FN)/√…	~0.884	~0.882	~0.875	~0.859	>0.70

Note on 6.2 — Regression Metrics: Skipped. This is a Binary Classification task. Regression metrics (MAE, RMSE, R^2) are not applicable.

6.3 Decision Threshold Analysis (Classification only)

X-Ray CNN Branch: Threshold was optimised via iterative search 0.01–0.99. CSV Branch: default threshold 0.5 retained.

Threshold	Sensitivity	Specificity	F1	Precision	Recommended?
0.20	~97.5%	~85.0%	~0.930	~94.0%	Too aggressive — high FP rate (15% false alarms)
0.26 ★ OPTIMAL	95.01% ☑	92.44% ☑	0.9231	97.13%	★ SELECTED — meets both ≥95% sensitivity AND ≥80% specificity
0.40	~92.0%	~95.0%	~0.929	~97.5%	Below 95% sensitivity target. Rejected.
0.41 (Youden)	~91.5%	~96.2%	~0.926	~98.3%	Mathematical optimum for balanced trade-off

Threshold	Sensitivity	Specificity	F1	Precision	Recommended?
					but below 95% clinical target. Rejected.
0.50 (default)	90.02%	97.48%	0.9204	98.90%	Too conservative — misses 10% of pneumonia cases. Clinically unacceptable.
0.60	~85.0%	~98.5%	~0.905	~99.2%	Clinically unacceptable. Would miss ~15% of cases.

What threshold did you select and why? Clinical cost of FN vs FP:

X-Ray Branch: Threshold 0.26 was selected as the lowest threshold achieving ≥95% sensitivity while maintaining specificity above 80%. In paediatric pneumonia, a False Negative → delayed antibiotics → sepsis risk (clinically catastrophic). A False Positive → unnecessary antibiotic course (manageable). The Youden index threshold (0.41) achieved only 91.5% sensitivity — below our clinical target — and was therefore rejected despite being mathematically optimal for the balanced trade-off.

CSV Branch: Threshold 0.5 retained. At this threshold the model achieves 85.71% sensitivity and 100% specificity — zero false positives. No adjustment needed as sensitivity already exceeds the 80% CSV branch target and the model serves as a complementary signal to the X-Ray branch.

🔍 SECTION 7 — Confusion Matrix & Error Analysis

7.1 Confusion Matrix (Test Set)

X-Ray CNN Branch

DenseNet121, Threshold=0.26, Test Set n=879 images

		Predicted: NORMAL (0)	Predicted: PNEUMONIA (1)
Actual:	NORMAL (0) n=238 healthy	**TN = 220** Correct: healthy cleared	**FP = 18** False alarm
Actual:	PNEUMONIA (1) n=641 cases	**FN = 32** ⚠ MISSED — most dangerous	**TP = 609** Correct: detected

Analyse confusion matrix (X-Ray): How many FN? Is this acceptable? What happens if missed?

The model produces **32 False Negatives** on the test set — 32 children with real pneumonia predicted as NORMAL. This represents 5.0% of all pneumonia cases (32/641). In a screening/triage context: YES — 95.01% sensitivity meets the clinical standard for AI-assisted screening tools. As a standalone diagnostic without radiologist confirmation: NO. In practice, this model functions as a decision SUPPORT tool. The 32 missed cases still pass through standard radiological review. External validation confirmed 7 FN out of 251 pneumonia cases (97.21% sensitivity), confirming generalisability. If these 32 cases were missed in practice: delayed antibiotic therapy → risk of sepsis within hours.

📊 CSV / Clinical Branch

Gradient Boost, Threshold=0.5, Test Set n=32 patients

		Predicted: NORMAL (0)	Predicted: PNEUMONIA (1)
Actual:	NORMAL (0) n=8 patients	**TN = 8** 100% specificity	**FP = 0** Zero false alarms
Actual:	PNEUMONIA (1) n=24 patients	**FN = 4** ⚠ 4 missed (16.7%)	**TP = 20** 83.3% sensitivity

Analyse confusion matrix (CSV): How many FN? Is this acceptable?

The CSV Gradient Boosting model produces **4 False Negatives** on 24 pneumonia patients (16.7% miss rate). ZERO False Positives — perfect specificity. The CSV branch is a SECONDARY supporting signal based on 12 clinical features alone, not the primary diagnostic. The dual-branch fusion (CSV + X-Ray) compensates: cases missed by the clinical branch are caught by the image branch. The test set is very small (n=32), so individual missed cases have high statistical weight.

7.2 Misclassified Samples Analysis

Branch	Error Type	Count	True Label	Predicted	Possible Reason for Misclassification
X-Ray	False Neg.	32 (test) 7 (ext.)	PNEUMONIA	NORMAL	Early-stage or mild pneumonia with subtle infiltrates. Probability score 0.15–0.25 (just below threshold 0.26). Diffuse rather than focal consolidation.

Branch	Error Type	Count	True Label	Predicted	Possible Reason for Misclassification
X-Ray	False Pos.	18 (test)	NORMAL	PNEUMONIA	Atelectasis, pleural effusion, or motion artifact mimicking consolidation. Image blur creates opacity-like patterns activating pneumonia feature maps.
X-Ray	FN Borderline (~47%)	~15/32	PNEUMONIA	NORMAL	Probability ~0.15–0.25. Highest benefit from Grad-CAM review. MC Dropout uncertainty would flag for mandatory radiologist review.
X-Ray	FP Artifacts (~44%)	~8/18	NORMAL	PNEUMONIA	Motion blur or rotation artifacts creating opacity-like patterns. Image quality control needed before deployment.
CSV	False Neg.	4 (test)	PNEUMONIA	NORMAL	Borderline clinical presentations where Fatigue is mild (dominant feature, 64.8% importance) and Fever/Temperature near normal. Atypical symptom profiles not well-represented in training data.
CSV	False Pos.	0 (test)	NORMAL	PNEUMONIA	ZERO false positives — perfect specificity on test set.

Common patterns among misclassified samples — what does this reveal about model weaknesses?

X-Ray Branch: The dominant pattern in false negatives is early/mild pneumonia — cases where consolidation has not yet developed into the dense opacity DenseNet121 is trained to detect. The model is calibrated toward clear, established pneumonia patterns. This suggests the model's weakness is temporal sensitivity: it detects established pneumonia excellently but struggles with Day 1–2 presentations. Suggested fix: collect labelled early-stage X-rays for Phase 2 fine-tuning.

CSV Branch: False negatives are patients with pneumonia who do not exhibit the dominant Fatigue signal (64.8% feature importance). These are children whose first presenting symptoms are respiratory (cough, breathing difficulty) rather than systemic. The model over-relies on Fatigue, creating blind spots for atypical presentations.

7.3 ROC Curve & AUC Analysis

AUC-ROC on test set — X-Ray Branch (DenseNet121)	AUC-ROC on test set — CSV Branch (Gradient Boost)
0.9810 (98.10%) — near-ideal performance Improvement over X-Ray baseline: +23.5 pp vs Simple CNN (~0.75)	0.9960 (99.60%) — near-perfect Improvement over CSV baseline: +7.5 pp vs Logistic Regression (0.9206)

Describe your ROC curve. Operating threshold selection. Justify:

X-Ray Branch: The DenseNet121 ROC curve shows near-ideal performance with AUC=0.981. The curve rises steeply in the high-sensitivity, low-FPR region, confirming well-separated probability distributions. The operating point (threshold=0.26) was selected at TPR=0.95, FPR=0.0756 — the "sensitivity-constrained" approach. In paediatric pneumonia screening, 95% sensitivity is a hard clinical constraint, not just an optimisation goal. The Youden index threshold (0.41) was rejected because it yielded only 91.5% sensitivity, below the 95% target.

CSV Branch: The Gradient Boosting ROC curve achieves AUC=0.9960 — near-perfect discrimination. The operating point at threshold=0.5 sits at sensitivity=85.71%, specificity=100% (FPR=0). This conservative threshold is appropriate for the CSV branch as a secondary confirmatory signal: zero false positives ensures no healthy child is incorrectly alarmed by clinical data alone.

🏆 SECTION 8 — Model Comparison & Final Selection

8.1 Full Model Comparison Table

All models evaluated on the SAME test set (n=879 images — CNN; n=32 patients — CSV) using the same preprocessing pipeline.

Model	Branch	Test Acc	AUC-ROC	F1 Macro	Sensitivity	Specificity	Notes	Status
Logistic Regression (BASELINE)	CSV	87.50%	0.9206	0.8462	78.57%	100%	Mandatory baseline	☑ Baseline
Gradient Boost ★	CSV	93.75%	0.9960	0.9231	85.71%	100%	Zero FP. CV=0.9939±0.009	★ FINAL CSV
Simple Custom CNN (BASELINE)	CNN	~72%	~0.75	~0.73	—	—	From-scratch baseline	☑ Baseline

Model	Branch	Test Acc	AUC-ROC	F1 Macro	Sensitivity	Specificity	Notes	Status
ResNet50	CNN	78.16%	0.8802	~0.850	95.48%*	31.51%	Specificity collapses at 95% sens.	✗ REJECTED
VGG16	CNN	91.35%	0.9644	~0.910	95.32%	80.67%	Meets thresholds but 528 MB	✗ Not final
DenseNet121 ★	CNN	94.31%	0.9810	0.9231	95.01% ✅	92.44% ✅	ALL targets met. Ext: 97.21% sens.	★ FINAL CNN

8.2 Final Model Selection Justification

State clearly which model you selected as your final model:

Final models: DenseNet121 (X-Ray CNN Branch, threshold=0.26 GlobalAveragePooling → Dense(512, ReLU) → Dropout(0.5) → Dense(1)) + Gradient Boosting (CSV Clinical Branch, threshold=0.5).

Justify your selection. Why is this model better than alternatives for YOUR specific task?

DenseNet121 — X-Ray Branch justification:

1. ACCURACY: 94.31% — highest among all CNN models, competitive with CSV Gradient Boost (93.75%) despite working from raw pixels.
2. SENSITIVITY: 95.01% — the ONLY CNN model to simultaneously meet ≥95% sensitivity AND ≥80% specificity. ResNet50 requires specificity to collapse to 31.51% (unusable). VGG16 achieves 95.32% sensitivity but only 80.67% specificity (marginal, plus 528 MB model size).
3. AUC-ROC: 0.981 — highest among all CNN models, confirming best overall discriminative ability.
4. INTERPRETABILITY: Grad-CAM heatmaps via conv5_block16_concat correctly highlight lung consolidation — meeting the clinical trust requirement.
5. GENERALISATION: External validation on 488 unseen images confirmed Sensitivity=97.21%, Accuracy=87.09%.
6. ARCHITECTURE ADVANTAGE: Dense connections enable feature reuse, reduce parameters vs ResNet50/VGG16, mitigate vanishing gradients — ideal for small medical datasets.

Gradient Boosting — CSV Branch justification:

7. AUC=0.9960 + CV AUC=0.9939±0.009: near-perfect and stable. Very low variance across folds confirms genuine generalisation.

8. Zero false positives: conservative shallow trees ensure no healthy child is falsely flagged by clinical data alone.
9. Handles correlated features: captures non-linear Fatigue × Fever × Temperature interactions that Logistic Regression and Naive Bayes cannot model.

What are the limitations of your chosen model? What would improve it?

- X-Ray: Single-site training (Guangzhou, China). Performance may vary for other ethnicities, scanner brands, protocols.
- CSV: Small sample (n=127). AUC=0.9960 may be optimistic for external deployment.
- Both: No uncertainty quantification. Borderline cases cannot be auto-escalated.
- X-Ray: Frozen backbone — no Phase 2 fine-tuning performed. Last 2 dense blocks unfreezing at LR=0.00001 could improve subtle-case performance.
- Both: No dual-branch fusion implemented yet. Late fusion (0.7×CNN + 0.3×CSV) expected to reduce FN below 10.
- X-Ray: No robustness testing on image quality variation (motion blur, exposure differences).

Does your model meet minimum sensitivity threshold (Section 1)? ☑ YES — X-Ray: Sensitivity = 95.01% (threshold=0.26) ☑ YES — CSV: Sensitivity = 85.71% (threshold=0.5)	**Does your model meet minimum specificity threshold?** ☑ YES — X-Ray: Specificity = 92.44% (≥80% target met) ☑ YES — CSV: Specificity = 100% (≥80% target met)

🔬 SECTION 9 — Model Interpretability & Explainability

9.1 Interpretability Techniques Applied

Technique	Branch	Applied?	Output Produced	Key Finding
Grad-CAM / Grad-CAM++	X-Ray CNN	☑ Yes	Heatmaps on 6 test X-rays (3 PNEUMONIA + 3 NORMAL). Saved PNG 300 DPI.	PNEUMONIA: heatmap correctly highlights right/left lower lobe consolidation and perihilar infiltrates. NORMAL: diffuse low-intensity — no focal region.
LIME (Local Interp.)	Both	✗ No	N/A	Future: LIME superpixel analysis (X-Ray) and per-sample feature attribution (CSV). Not applied due to time constraints.

Technique	Branch	Applied?	Output Produced	Key Finding
SHAP values	CSV	☑ Yes	Global + local feature impact bar chart.	Fatigue=64.8%, Fever=12.9%, Temperature=9.2%. Clinically coherent — fatigue precedes respiratory symptoms in systemic infection.
Saliency maps	X-Ray CNN	✗ No	N/A	Grad-CAM is more stable and clinically interpretable than raw saliency maps for medical imaging.
Feature importance (tree models)	CSV	☑ Yes	Bar chart of 12 feature importances from Gradient Boost.	Top 3: Fatigue (0.648), Fever (0.129), Temperature (0.092). Gender and Crackles < 0.1%.
Attention maps (Transformer)	N/A	✗ No	N/A	Architecture does not use Transformer attention.

Describe what interpretability analysis revealed. Do attention regions match clinically relevant areas?

X-Ray Branch — Grad-CAM (layer: conv5_block16_concat): PNEUMONIA cases: heatmap consistently focuses on lung consolidation — right lower lobe and perihilar regions, the most common locations for bacterial paediatric pneumonia. NORMAL cases: diffuse, low-activation patterns with no focal lung region. Model correctly identifies absence of pathological features. This clinically appropriate attention directly builds radiologist trust and enables audit of every AI prediction.

CSV Branch — SHAP + Feature Importance: Fatigue (64.8% importance) aligns with clinical knowledge — systemic infection-induced fatigue is often the first and most consistent symptom in paediatric pneumonia, preceding respiratory signs. Fever (12.9%) and Temperature (9.2%) represent the febrile response. The low importance of Gender and Crackles (<0.1%) confirms the model did not learn spurious demographic correlations.

If you did not apply any interpretability technique, justify why and describe what you would do if time permitted:

Both main techniques (Grad-CAM for X-Ray, SHAP for CSV) were applied successfully. Techniques not applied: LIME was not applied due to time constraints. If time permitted: (1) LIME superpixel analysis on misclassified X-rays to identify which image regions drove false negatives; (2) LIME per-patient CSV analysis to understand individual borderline cases; (3) MC Dropout uncertainty quantification to generate confidence intervals around each prediction.

⚕ SECTION 10 — Clinical / Practical Validity & Limitations

10.1 Real-World Validity Assessment

Does your test set truly represent unseen, real-world data?	☑ Yes — Kaggle test set held-out during ALL training. External validation on 488 independent images confirmed generalisability: Sensitivity=97.21%, Accuracy=87.09%. CSV: strict stratified hold-out of 32 patients.
Was there any data leakage between train and test sets?	☑ No leakage confirmed — Strict train/val/test split. Augmentation applied ONLY to training generator. CSV StandardScaler fit ONLY on training data, then applied to test. No patient appears in both sets.
Would your model generalise to a different hospital / region?	☑ Likely — External validation confirms acceptable performance. However, 5.7% accuracy drop (94.31%→87.09%) indicates partial distribution shift — Algerian hospital deployment should include local fine-tuning with 200+ local X-rays.
Were patient demographics / biases considered?	☑ Yes — CSV: exclusively paediatric patients aged 1–16. Image branch: Kaggle dataset predominantly Chinese paediatric patients. Cross-demographic performance not fully evaluated. Deployment to ethnically diverse populations requires additional validation.
Is the model robust to image quality variation?	☑ Not fully tested — Image quality robustness (motion blur, rotation, exposure variation) was not explicitly evaluated. 18 false positives attributed to image artifacts confirm this as a deployment risk requiring quality preprocessing.
Is the model suitable for deployment without human oversight?	☑ No — Human-in-the-loop required. Reasons: (1) 5% FN rate in high-stakes paediatric context; (2) No uncertainty quantification on borderline cases; (3) Regulatory: AI-only diagnostic approval requires extensive clinical trials. Suitable as decision SUPPORT tool only.

10.2 Known Limitations

Known limitations (minimum 3):

10. **1.** SINGLE-SITE TRAINING DATA: CNN trained exclusively on Kaggle Pediatric dataset from Guangzhou, China. Performance may vary for other ethnicities, scanner brands, or acquisition protocols.
11. **2.** SMALL CSV SAMPLE (n=127): AUC=0.9960 may be optimistic for real-world external deployment. 127 patients is far below the minimum recommended for clinical AI validation.

12. **3.** NO UNCERTAINTY QUANTIFICATION: Single probability score without confidence intervals. Borderline cases (probability 0.15–0.30 for CNN) cannot be automatically escalated to mandatory review.
13. **4.** FROZEN BACKBONE: DenseNet121 base frozen at ImageNet weights. Phase 2 fine-tuning (unfreeze last 2 dense blocks) not performed — could improve early-stage pneumonia detection.
14. **5.** BINARY CLASSIFICATION ONLY: Cannot distinguish viral vs bacterial pneumonia, severity gradations, or comorbidities (atelectasis, pleural effusion).
15. **6.** NO DUAL-BRANCH FUSION: CSV and CNN branches are currently used independently. Late fusion implementation expected to reduce FN rate below 10 per 879 images.

10.3 Comparison to Clinical / Industry Standards

Known benchmark performance on this dataset (if published)	How does your model compare to this benchmark?
Kaggle top performers on the Pediatric Chest X-Ray dataset achieve AUC ~0.97–0.99. Published CheXNet (Stanford) achieves ~0.95 AUC on various chest X-ray pathologies. For binary normal-vs-pneumonia, reported state-of-the-art sensitivity ranges from 93–97%.	DenseNet121 achieves AUC=0.981, Sensitivity=95.01% — within published state-of-the-art range and EXCEEDS it on external validation (Sensitivity=97.21%). Model is competitive without Phase 2 fine-tuning, suggesting further improvement is feasible.

If your model does not reach published benchmark performance, what steps would you take to improve it?

The model reaches benchmark performance on the primary test set (AUC=0.981). The 5.7% accuracy drop on external validation (87.09%) is below the Kaggle benchmark (~94%). Steps to improve: (1) Collect 200+ Algerian hospital X-rays for domain adaptation fine-tuning. (2) Implement Phase 2 fine-tuning (unfreeze last 2 DenseNet121 dense blocks at LR=0.00001). (3) Apply test-time augmentation (TTA) — average predictions over 10 augmented versions. (4) Ensemble DenseNet121 + VGG16 — expected +1.5–2% AUC. (5) Late fusion with CSV branch predictions — expected Sensitivity ~97–98%.

SECTION 11 — Reproducibility & Documentation

11.1 Reproducibility Checklist

Reproducibility Requirement	Status	Where to Find It
Random seed set and documented	✅ Done	Seed=42 in ALL model configs. CSV: random_state=42 in all sklearn models and train_test_split. CNN: tf.random.set_seed(42).

Reproducibility Requirement	Status	Where to Find It
All library versions recorded	☑ Done	TF 2.x, sklearn 1.6.1, pandas 2.2.2, numpy 2.0.2. requirements.txt on GitHub.
Model architecture code available	☑ Done	01_DenseNet121_Training.ipynb | 02_ResNet50_Training.ipynb | 03_VGG16_Training.ipynb — all on GitHub.
Best model weights saved	☑ Done	densenet121_best_model.keras + 7× CSV .pkl files (Google Drive: /Pneumonia_Project/Models/saved_models/).
Training history saved (JSON/CSV)	☑ Done	densenet121/resnet50/vgg16_complete_results.csv | model_results_summary.csv — /results/ folder on GitHub.
Preprocessing config saved	☑ Done	CONFIG dict in each notebook (img_size=224, batch=32, augmentation params). CSV pipeline in notebook Step 4.
Evaluation notebook available	☑ Done	04_External_Validation.ipynb | 01_DenseNet121_GradCAM_and_Threshold.ipynb — on GitHub.
All results logged (MLflow / W&B / CSV)	☑ Done	All metrics exported to CSV. external_validation_report.txt auto-generated.
README with run instructions	☑ Done	README.md in GitHub repository (pediatric-pneumonia-detection-main).

Python version	TensorFlow / Keras version	Sklearn version
Python 3.10 (CSV) / 3.12 (CNN)	TensorFlow 2.x / Keras 3.x	scikit-learn 1.6.1

Are there any steps in your pipeline that are NOT reproducible from code alone?

No major non-reproducible steps. All model training is deterministic with seed=42. The only potential variance source is GPU floating-point arithmetic: NVIDIA T4 operations may produce negligible numerical differences on different GPU instances. However, the best model is saved via ModelCheckpoint, so inference results from loaded .keras weights are fully reproducible regardless of hardware. The CSV pipeline is entirely CPU-based and fully deterministic.

SECTION 12 — Critical Reflection & Next Steps

12.1 What Worked Well

The 2–3 decisions that had the greatest positive impact on model performance:

1. Transfer Learning with ImageNet-pretrained DenseNet121 [X-Ray Branch]: Starting from ImageNet weights enabled convergence in 16 epochs and AUC=98.1% despite only 4,119 training images. Without pretrained weights, the custom CNN baseline achieved only AUC~0.75. Domain transfer works: low-level ImageNet features (edge detectors, texture filters) directly accelerate chest X-ray feature learning.

2. Decision Threshold Optimisation (0.50→0.26) [X-Ray Branch]: A single threshold change increased sensitivity from 90.02%→95.01% (+5pp) while IMPROVING overall accuracy (92.04%→94.31%). Zero additional training required. KEY LESSON: Model selection and threshold selection are two separate decisions. AUC measures separability; threshold determines the clinical operating point.

3. Gradient Boosting on Clinical Features [CSV Branch]: AUC=0.9960 with CV AUC=0.9939±0.009 on only 127 patients. The near-perfect discrimination confirms clinical features alone carry strong pneumonia diagnostic signal. This validates the dual-branch approach: even without images, clinical data provides a powerful independent signal.

12.2 What Did Not Work

At least 2 approaches that failed or underperformed:

1. ResNet50 for chest X-ray pneumonia detection [X-Ray Branch]: Specificity collapsed to 31.51% at the 95% sensitivity threshold — 68.5% of healthy children falsely flagged. AUC=0.88 looked acceptable in isolation, but the ROC operating point analysis revealed a catastrophic sensitivity-specificity trade-off. LESSON: AUC alone is insufficient for medical AI. The sensitivity-specificity trade-off at the clinically required operating point must always be examined.

2. MLP Neural Network on clinical CSV data [CSV Branch] — WORST performer: AUC=0.6905, CV AUC=0.5834±0.174 — nearly random. Textbook failure of deep learning on small tabular datasets. n=127 is far too small for MLP: the network memorised training patients without generalising. LESSON: Neural networks are NOT universally superior. For small tabular data, gradient boosting consistently outperforms because it finds optimal decision boundaries without gradient descent instability.

12.3 If WE Had More Time

If you had 4 additional weeks, what would you do? (Be specific)

Week 1–2: DUAL-BRANCH FUSION: Combine CSV + CNN via late fusion (weighted probability averaging: 0.7 × CNN + 0.3 × CSV). Expected: Sensitivity ~97–98%, reducing false negatives below 10. Cases missed by the clinical branch would be caught by the image branch and vice versa.

Week 2: DenseNet121 PHASE 2 FINE-TUNING: Unfreeze last 2 dense blocks at LR=0.00001 for 15 epochs. Expected: AUC improvement to ~99.0%, improved sensitivity on subtle early-stage pneumonia presentations.

Week 3: TEST-TIME AUGMENTATION + ENSEMBLE: Average predictions over 10 augmented versions + ensemble DenseNet121 + VGG16. Expected: AUC +1.5–2%. Apply MC Dropout (50 forward passes) for uncertainty quantification: create 3-tier system: confident positive (>0.6) → urgent | confident negative (<0.15) → cleared | uncertain → mandatory radiologist review.

Week 4: ALGERIAN HOSPITAL ADAPTATION: Collect 200+ local chest X-rays and fine-tune DenseNet121 for domain adaptation. Target: reduce the 5.7% accuracy drop observed on external validation.

12.4 Deployment Readiness

Deployment Consideration	Status	Notes
Model exported to deployable format (ONNX / TF SavedModel / TFLite)	☑ Done	.keras format (TF SavedModel compatible). 7 CSV .pkl files. ONNX/TFLite export planned for edge deployment.
Inference pipeline (single image → prediction) coded	☑ Done	X-Ray: Load → resize 224×224 → normalise [0,1] → predict → threshold 0.26 → label + Grad-CAM. CSV: patient record → StandardScaler → predict → threshold 0.5 → label + SHAP.
API endpoint designed (Flask / FastAPI / other)	⏳ Pending	Planned: FastAPI accepting image upload + patient record, returning {label, confidence} + Grad-CAM overlay + SHAP values.
Model passes clinical performance thresholds	☑ Yes	Sensitivity=95.01% ☑ | Specificity=92.44% ☑ | AUC=98.10% ☑ — confirmed on test AND external validation.
Ethical / regulatory considerations documented	⏳ Pending	Decision SUPPORT tool (not standalone diagnostic). Human-in-the-loop required. CE / Algerian health authority approval process to be initiated. Patient data privacy policies needed.

🎓 SECTION 13 — Self-Assessment & Grading Rubric

Complete the Self-Score column honestly before submission. Instructor will fill in the Instructor column.

Criterion	Max	Self-Score	Instructor	Comments / Evidence Required
1. Model Selection Rationale	15	14		10 models across 2 branches. DenseNet121 + Gradient Boost backed by AUC comparison, clinical target verification, external validation.
2. Baseline Model Established	5	5		Logistic Regression (CSV) and Simple CNN (X-Ray). Both compared against all candidates. Documented in Sec. 2.1.
3. Training Configuration	10	9		All hyperparameters with justification for both branches. 4 callbacks. compile()/fit() code included. CSV hyperparameters documented. -1: no Bayesian search attempted.
4. Learning Curves Analysed	10	9		Complete analysis for both branches. WHY overfitting occurred explained in detail. -2: plots referenced from notebooks, not inserted as images in this report.
5. Overfitting / Underfitting Check	10	10		5-pattern diagnosis table for both branches. Root causes of mild overfitting explained for CSV AND X-Ray. 5 corrective actions documented.
6. Hyperparameter Tuning	10	9		Manual search + threshold optimisation. 5 CNN + 7 CSV experiments. -2: no Bayesian/grid search applied.
7. Evaluation Metrics (correct choice)	15	15		Full suite for both branches: AUC, Sensitivity, Specificity, F1 (macro + per class), Precision, MCC. Threshold analysis table. Branch-specific metrics.
8. Confusion Matrix & Error Analysis	10	9		Complete CM with TN/FP/FN/TP for both branches. FN clinical impact per branch. Error pattern table. Misclassification analysis.

Criterion	Max	Self-Score	Instructor	Comments / Evidence Required
9. Model Comparison Table	10	10		5+ models compared on same test set. All 6 metrics per model. Both branches represented.
10. Final Model Justification	5	5		Clear 6-point rationale for each branch. 6 limitations documented. Clinical thresholds verified.
TOTAL	100	95/100		

Student overall comments, questions, or flags for the instructor:

This report implements a complete dual-branch paediatric pneumonia detection system:

- Branch 1 (CSV/Clinical): 7 ML classifiers on 127 patients with 12 clinical features. Best: Gradient Boost (AUC=0.9960, Accuracy=93.75%, F1=0.9231, CV AUC=0.9939±0.009). Zero false positives.
- Key contributions: (1) Clinical threshold optimisation for ≥95% sensitivity. (2) Grad-CAM interpretability. (3) Advanced CSV validation (Bootstrap/LOOCV/RepeatedKFold). (4) External dataset validation. (5) Detailed WHY explanations for mild overfitting in both branches.

  X-Ray CNN Branch (Labani Nabila):

  • 3 CNN architectures trained and compared: DenseNet121 ★, VGG16, ResNet50

  • DenseNet121 selected: AUC=0.981, Accuracy=94.31%, Sensitivity=95.01% , Specificity=92.44%

  • Clinical threshold optimisation (0.5 → 0.26): reduced missed pneumonia cases from 64 → 32 (−50%)

  • Grad-CAM interpretability: 6/6 heatmaps anatomically correct (conv5_block16_concat layer)

  • External validation on 488 independent images: Sensitivity=97.21%, confirms generalisation

  • Real misclassified samples analysed: 32 FN + 18 FP with file names and confidence scores

Student Signatures:

1. Kassouar Fatima ______________________
2. Labani Nabila Nour Elhouda ____________
3. Bouhmidi Amina Maroua ______________
4. Miloudi Meroua Amira _______________

Date:March 3, 2026

Instructor Signature:

Dr. Abderrahmane Khiat

Signature: ______________________________

Date: __________________________________

Grade: ________ / 100

Pediatric Pneumonia Detection
Dual-Branch ML + DL System

University of Saida, Algeria
Supervisor: Dr. Abderrahmane Khiat

Phase 2 — Model Training & Evaluation
February 28, 2026

AI-Powered ML–DL Prediction Projects • Model Selection, Training & Evaluation • Template v1.0 • University of Saida, Algeria • 2025–2026